

The Computing Revolution in Biosciences

BU-ISCIII

Unidades Comunes Científico Técnicas - SGSAFI-ISCIII

2428 Octubre 2022, 4ª Edición
Programa Formación Continua, ISCIII

Index

The Computing Revolution in Biosciences:

- The Century of Biology
- Computing in Biosciences
- The Omics Era
- Change of Paradigm
- HPC infrastructure
- Workflows
- The need of standardisation
- Nextflow
- Containers
- Singularity

The Century of Biology

“If the 20th century was the century of physics, the 21st century will be the century of biology. While combustion, electricity and nuclear power defined scientific advance in the last century, the new biology of genome research—which will provide the complete genetic blueprint of a species, including the human species—will define the next.”

VENTER, C., & COHEN, D. (2004). The Century of Biology. *New Perspectives Quarterly*, 21(4), 73–77.
doi:10.1111/j.1540-5842.2004.00701.x

Healthcare IT News GLOBAL EDITION

Obama's next move: Precision medicine and genomics venture capitalist?

By [Jessica Davis](#) | June 29, 2016 | 04:48 PM



Healthcare IT News GLOBAL EDITION

Microsoft, Google invest in precision medicine startup DNAnexus

By [Bernie Monegain](#) | January 02, 2018 | 12:25 PM



Computing in Biosciences I

Research used to be focussed in a small number of samples and researchers analysed them with the whatever means they had and/or felt more comfortable with:

- Windows based PC using programs with visual interface
- Macs and Linux based workstations
- Remote web servers
- Web-based platforms (i.e. Galaxy) and remote HPC
- HPC local environments

Computing in Biosciences II

- Windows based PC using programs with visual interface

| Pros | Cons |
|--------------------------|--|
| Data remains private | No backups or data management schemes |
| Software easy to install | Software version not easy to control, binaries are black boxes |
| Graphic interface | No control over hidden parameters |
| | Analysis are irreproducible |

Computing in Biosciences III

- Macs and Linux based workstations

| Pros | Cons |
|--|--|
| Data remains private | No backups or data management schemes |
| Control over software installed versions, open source programs | Software may not be easy to install, library and dependencies problems |
| All parameters are available for the command | Command line interface |
| | Analysis are irreproducible |

Computing in Biosciences IV

- Remote web servers

| Pros | Cons |
|---------------------------------------|--|
| No need to storage intermediate files | Your data is in someone else's computer No backups or data management schemes |
| No need to install software | Software version not easy to control, black boxes |
| Graphic interface | No control over hidden parameters |
| | Quotas Analysis are irreproducible |

Computing in Biosciences V

- Web-based platforms (i.e. Galaxy) and remote HPC

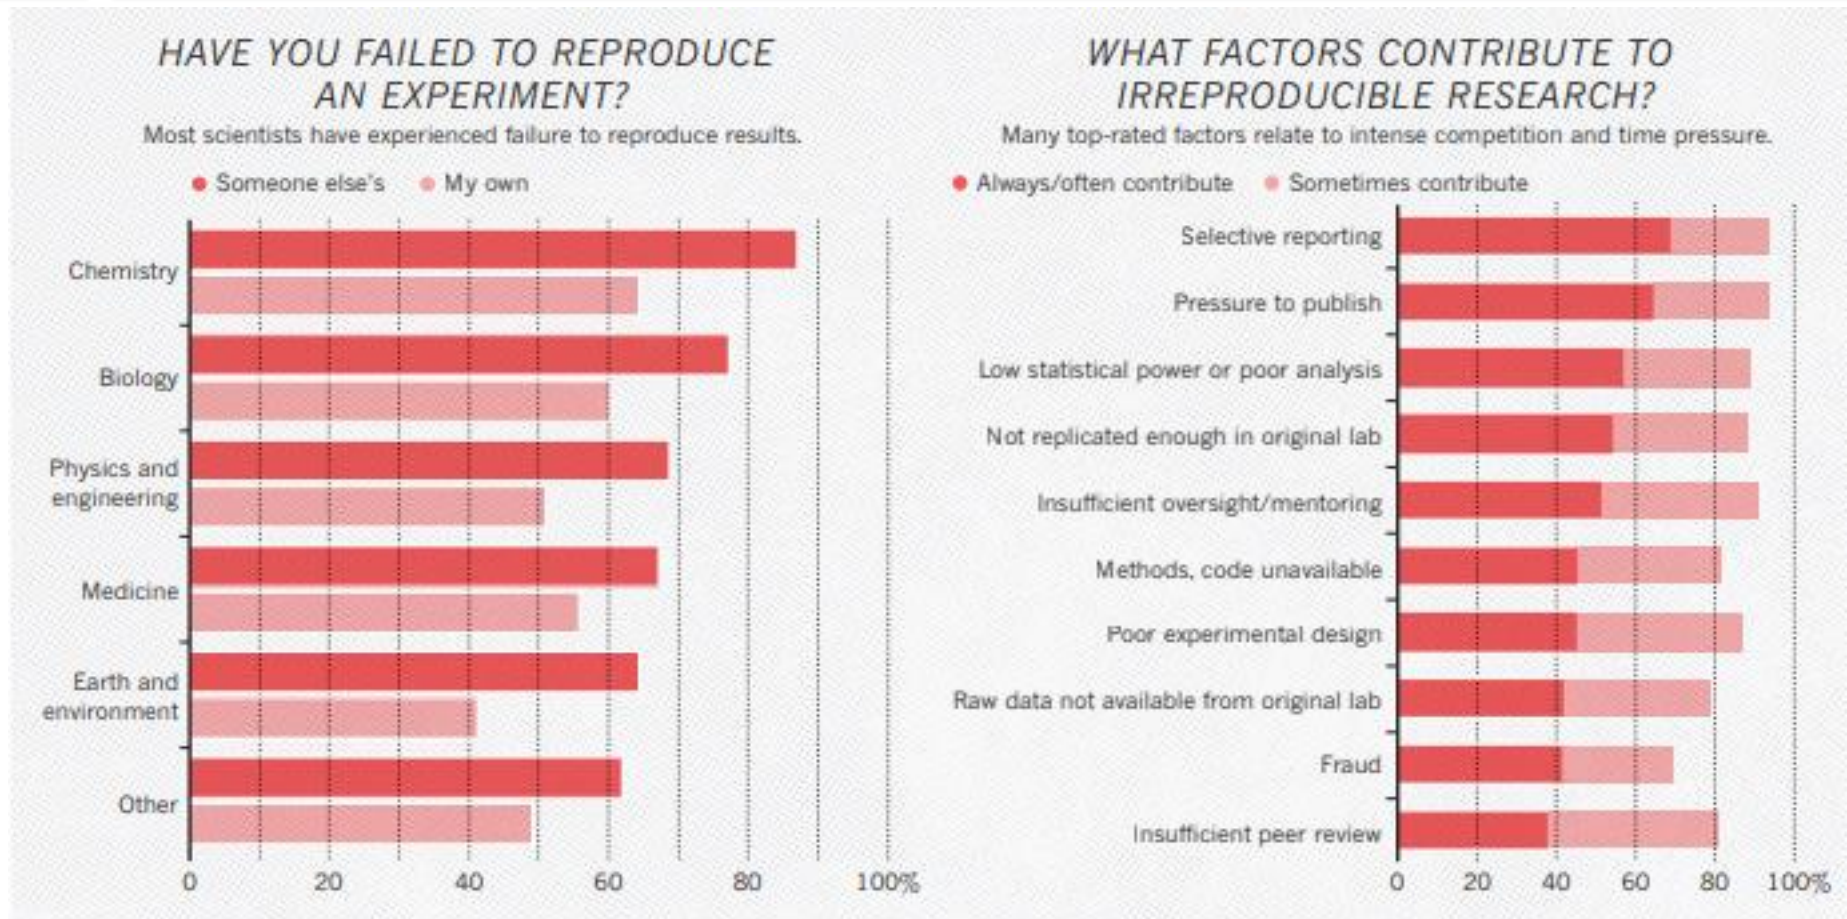
| Pros | Cons |
|--|--|
| No need to storage intermediate files | Your data is in someone else's computer No backups or data management schemes |
| No need to install software Partial control over installed software | No control over installed software, versions and future availability |
| Graphic interface | No control over hidden parameters |
| Analysis are partially reproducible | Quotas |

Computing in Biosciences VI

- HPC local environments

| Pros | Cons |
|--|--|
| Data remains private Backups and data management schemes | Quotas |
| No need to install software Partial control over installed software | No control over installed software, versions and future availability |
| All parameters are available for the command | Command line interface |
| Possibility of suggesting new software installations | Analysis may be irreproducible |

Is there a reproducibility crisis?



Source: Baker, M. "Reproducibility Crisis (Nature)," 3–5. doi:10.1038/533452A.

The Omics Era I


| | Coverage | No. of Reads | Read Length | BAM File Size | Strand NGS Size |
|--------------|----------|---------------|-------------|---------------|-----------------|
| Whole Genome | 37.7x | 975,000,000 | 115 | 82 GB | 104 GB |
| Whole Genome | 38.4x | 3,200,000,000 | 36 | 138 GB | 193 GB |
| Exome | 40x | 110,000,000 | 75 | 5.7 GB | 7.1 GB |

| Whole Genome Samples | Exome Samples | Space | Space including Backup |
|----------------------|---------------|--------|------------------------|
| 0 | 200 | 1.6 TB | 3.2 TB |
| 0 | 1000 | 8.0 TB | 16 TB |
| 100 | 0 | 15 TB | 30 TB |
| 1000 | 0 | 150 TB | 300 TB |
| 100 | 1000 | 23 TB | 46 TB |

Source: <https://www.strand-ngs.com/support/ngs-data-storage-requirements>

The Omics Era II

Total disk storage at EMBL-EBI



Installed (2008–2015) storage at EMBL-EBI. These figures include all installed storage, counting multiple backups for all data resources as well as unused storage to handle submissions in the immediate future

Source: Cook, Charles E et al. "The European Bioinformatics Institute in 2016: Data growth and integration" *Nucleic acids research* vol. 44,D1 (2015): D20-6.

Change of Paradigm I

1 sample

Research only: NGS was still a new thing, no applications 10 years ago

Reproducibility is not needed: Why would anyone reanalyse this?

Storage is not an issue: files of 1 sample fits everywhere in my HDD, maybe I will copy it in a CD-ROM

Computing is simple: no need to worry about resources or optimisation

multiple samples

Many applications: research, clinical, industrial, forensic, military, ...

Reproducibility, scalability , portability and standardisation are required

Storage is challenging: storage, indexation and backup required, privacy and legal standards

Computing requires optimisation and lots of resources

Change of Paradigm II

- Nowadays scientific computing paradigm

| Pros | Cons |
|---|--|
| Data remains private Backups and data management schemes | High storage space Dedicated file systems Databases to index files |
| Control over software installed versions, open source programs | Many versions of the same software coexists |
| All parameters are available for the command | You have to understand all software variations |
| Analysis are reproducible and public | You have to publish and document your work |


HPC infrastructure I

| Machine | OS | Software | CPU | RAM | Storage |
|-----------------------|------------|----------|-----------|--------|---------|
| Workstation (x5) | Centos 6.9 | /opt(*) | 4 cores | 32 Gb | 4 TB |
| Bioinfo01 (1 node) | | /opt(*) | 16 cores | 120 Gb | 500 Gb |
| HPC (16 nodes) | | /opt(*) | 320 cores | 8 TB | 500 Gb |

2 shared data storage disk boxes: 70TB + 250 TB

**VMs, ISCIII's Windows personal terminals, personal laptops mobile platforms,
cloud computing platforms, cloud storage, remote services, ...**

HPC infrastructure II



SYSTEM HIGHLIGHTS

The system will focus on features relevant to genomics researchers with features such as huge data storage capabilities, very high-memory research servers for maximum performance and integration with relevant biological databases.



7680 vCPU Cores

The CLIMB system is composed of over 7,500 CPU cores of processing power. This makes it the largest single system dedicated to Microbial Bioinformatics research, anywhere in the world.




500 Local Storage (TB)

To provide users with local, high performance, storage we have deployed IBM GPFS in each of the 4 sites, to provide 500TB of local storage. This storage is connected to our servers using Infiniband.



78 Total RAM (TB)

Unlike most supercomputers, the CLIMB system has been designed to provide large amounts of RAM, in order to meet the challenge of processing large, rich biological datasets. In comparison, the Spruce B supercomputer at the Atomic Weapons Research Establishment (number 68 on the Top 500 Supercomputers list, November 2014) has 35,000 cores, but only has 110 TB of RAM.




1000 Virtual Machines


The CLIMB system is not designed to provide a single HPC system, as is often the case within academic computing; rather, the CLIMB system provides a pool of CPU cores and RAM that Medical Microbial Bioinformatics researchers can gain access to. The system has been designed to support over 1,000 VMs running simultaneously, potentially supporting most of the Microbial Bioinformatics community within the UK.

Workflows I

- Bioinformatic analyses invariably involve shepherding files through a series of transformations, called a **pipeline** or a **workflow**.
- These transformations are done by executable **command line software** written for Unix-compatible operating systems.
- They need to be **reproducible**, **easy to maintain**, **portable** and **scalable**.



Workflows II



The need of standardisation I

Sequencing techniques are starting to be used in clinical diagnosis, and therefore workflows have to assure:

- **Reproducibility**

Results always have to be reproducible


- **Portability**

The analysis workflow must be executable in different platforms

- **Scalability**

The analysis workflow must be able to work with different numbers of samples

The need of standardisation II



Nextflow I

- **Nextflow** is a DSL for parallel and scalable computational pipelines.
- It enables **scalable and reproducible scientific workflows** using software **containers**.
- It **allows the adaptation of pipelines** written in the most common scripting languages.
- Its fluent **DSL** simplifies the implementation and the deployment of complex parallel and reactive workflows on clouds and clusters.

Nextflow II

Fast prototyping

Nextflow allows you to write a computational pipeline by making it simpler to put together many different tasks.

You may reuse your existing scripts and tools and you don't need to learn a new language or API to start using it.

Portable

Nextflow provides an abstraction layer between your pipeline's logic and the execution layer, so that it can be executed on multiple platforms without it changing.

It provides out of the box executors for SGE, LSF, SLURM, PBS and HTCondor batch schedulers and for [Kubernetes](#) and [Amazon AWS](#) cloud platforms.

Continuous checkpoints

All the intermediate results produced during the pipeline execution are automatically tracked.

This allows you to resume its execution, from the last successfully executed step, no matter what the reason was for it stopping.

Reproducibility

Nextflow supports [Docker](#) and [Singularity](#) containers technology.

This, along with the integration of the [GitHub](#) code sharing platform, allows you to write self-contained pipelines, manage versions and to rapidly reproduce any former configuration.

Unified parallelism

Nextflow is based on the *dataflow* programming model which greatly simplifies writing complex distributed pipelines.

Parallelisation is implicitly defined by the processes input and output declarations. The resulting applications are inherently parallel and can scale-up or scale-out, transparently, without having to adapt to a specific platform architecture.

Stream oriented

Nextflow extends the Unix pipes model with a fluent DSL, allowing you to handle complex stream interactions easily.

It promotes a programming approach, based on functional composition, that results in resilient and easily reproducible pipelines.

Nextflow III

Nextflow Report
Summary
Resources
Tasks
[elated_feynman]

Nextflow workflow report

[elated_feynman]

Workflow execution completed successfully!

Run times
Thu Aug 30 11:25:29 CEST 2018 - Fri Aug 31 12:46:56 CEST 2018 (completed a month ago, duration: **1d 1h 21m 27s**)

347 succeeded

Nextflow command

```
nextflow -C nextflow.config run /processing_Data/bioinformatics/pipelines/PikaVirus_nextflow/main.nf -with-report report -with-trace trace -with-timeline timeline -with-dag flowchart.png-resume
```

| | |
|-------------------|---|
| CPU-Hours | 152.8 |
| Launch directory | /processing_Data/bioinformatics/research/20180605_PIKAVIRUSNEXTFLOW_MJ |
| Work directory | /processing_Data/bioinformatics/research/20180605_PIKAVIRUSNEXTFLOW_MJ/work |
| Project directory | /processing_Data/bioinformatics/pipelines/PikaVirus_nextflow |
| Script name | main.nf |
| Script ID | 4212c5412541f99eb464fbb3b0ee39ae |
| Workflow session | dce5f6f6-6ed4-4ff9-a2f6-06f59b798d93 |
| Workflow profile | standard |
| Nextflow version | version 0.30.2, build 4867 (16-06-2018 17:49 UTC) |

Nextflow IV

Nextflow Report Summary Resources Tasks


[elated_feynman]

Resource Usage

These plots give an overview of the distribution of resource usage for each process.

CPU Usage

Raw Usage % Allocated



Memory Usage

Raw Usage % Allocated

Memory Usage

Nextflow V

Nextflow Report Summary Resources Tasks

[elated_feynman]

Tasks

This table shows information about each task in the workflow. Use the search box on the right to filter rows for specific values. Clicking headers will sort the table by that value and scrolling side to side will reveal more columns.

Values shown as: Human readable

Show 25 entries

Filter: Metrics Metadata All Search:

| task_id | process | tag | status | hash | allocated cpus | %cpu | allocated memory | %mem | vmem |
|---------|------------|----------|-----------|-----------|----------------|-------|------------------|------|---------|
| 1 | raw_fastqc | Sample11 | COMPLETED | e8/ece0ae | 1 | 97.6 | - | 0.1 | 1.8 GB |
| 2 | raw_fastqc | Sample10 | COMPLETED | ab/cdeea6 | 1 | 91.5 | - | 0.1 | 1.9 GB |
| 3 | trimming | Sample10 | COMPLETED | 4a/ea6dc1 | 1 | 612.3 | - | 1.6 | 36.5 GB |
| 4 | trimming | Sample11 | COMPLETED | 9e/c6fbd5 | 1 | 675.6 | - | 3.2 | 36.5 GB |

Containers I

Linux containers is a generic term for an implementation of operating system-level virtualization for the Linux operating system.

Containers allow us to **port** pipelines and **replicate** their exact execution environments across different hardware.

Currently, a number of such implementations exist, and they are all based on the **virtualization, isolation, and resource management mechanisms provided by the Linux kernel**.

Containers II

Singularity is a free, cross-platform and open-source computer program that performs operating-system-level virtualization.

One of the main uses of Singularity is to bring containers and **reproducibility to scientific computing** and the HPC world.

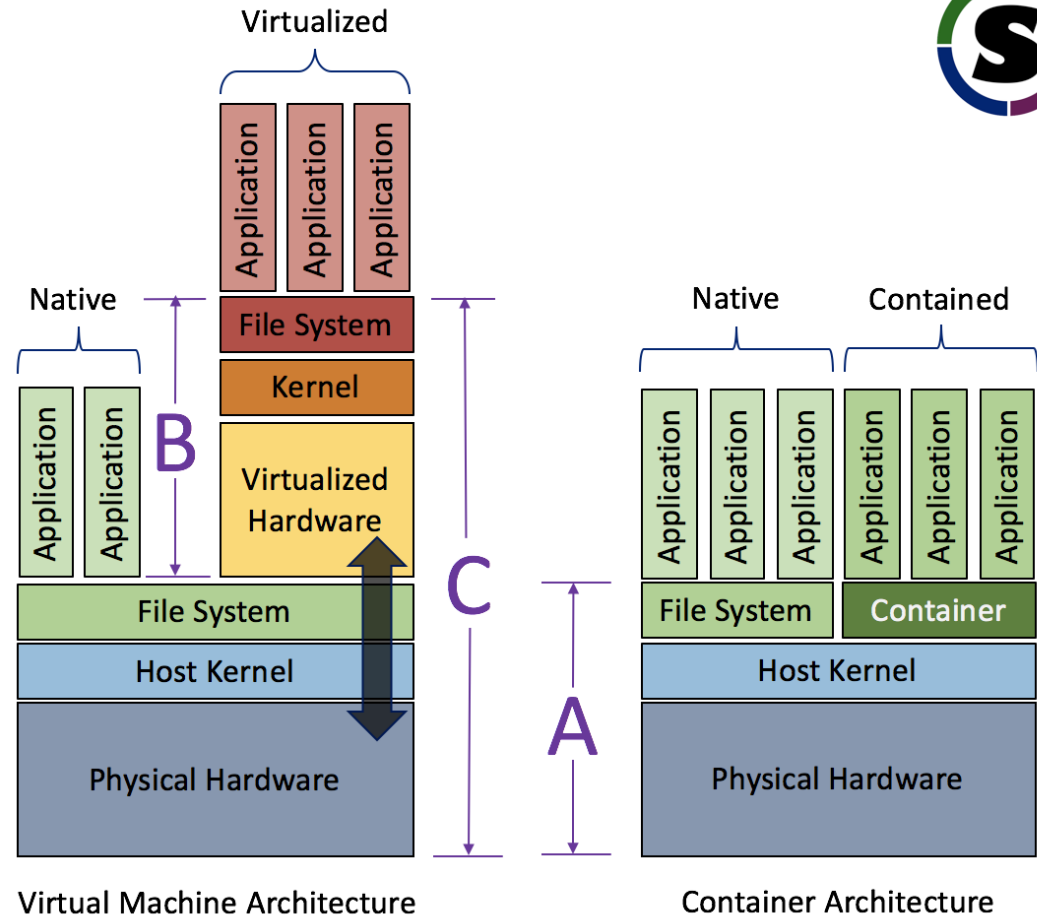
While Docker is broadly used, Singularity is **fully compatible with Docker**, plus Singularity does **not require root permissions** to be executed.

Singularity I



VIRTUAL MACHINES VS. CONTAINERS

- Applications running within a container will always be “closer” to the physical hardware
 - Notice how close to native a container behaves
- Applications running through a virtual machine will always have multiple levels of indirection
- The container’s proximity to the physical hardware equates to less overhead, higher performance and lower latency



Singularity II

Singularity image runs on the same level as the OS, directly above the kernel, and can access all hardware in the machine.

Not needing to virtualise the hardware and run a kernel again makes this kind of virtualisation really effective.

Filesystem is shared, and some paths are automatically mounted (/tmp and /home), while the others are optional.

Files of the host system can be created, modified and deleted from the image in the mounted folders.

Results I

Before the implementation of these combination of framework, software and guidelines, executing a workflow in an HPC environment consisted in the following steps:

- Loading input data
- Asking sysadmin to install dependencies
- Load references
- Estimate and book computational resources
- Manually execute each step of the pipeline, or automate it with a script
- Wait with no control over the process status until finished

Results II

Now a simple command works out of the box in any machine:

```
nextflow run //buisciii/main.nf -profile singularity
```

Plus, it give us:

- Dependency and computing automatization
- Resource usage statistics
- Easy to share and maintain
- Reproducibility and re-entrancy
- Transparency

Thanks for your attention!

And this is only the tip of the iceberg...
Check this if you wanna know what's really going under
the hood:



<https://github.com/BU-ISCIII>